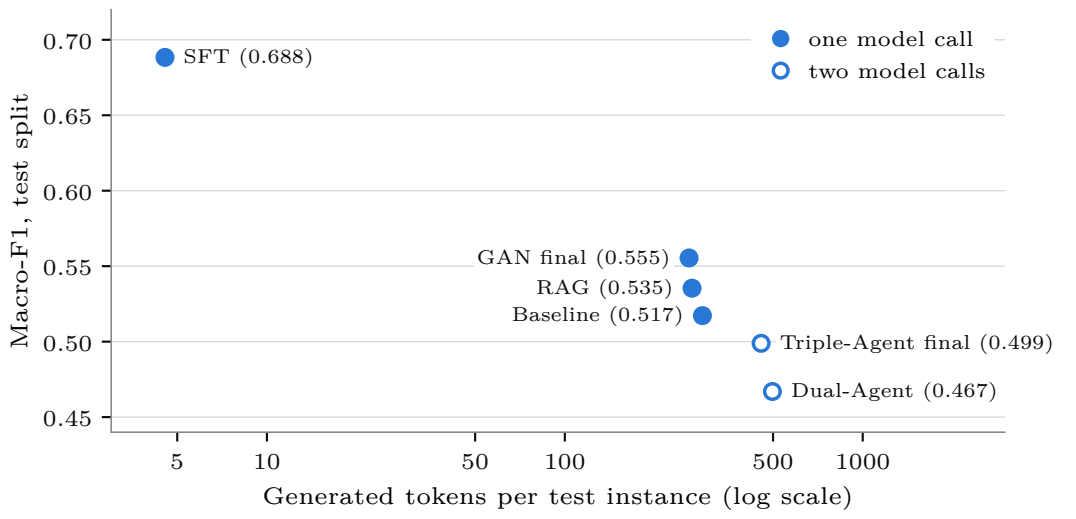
0.70
0.65
0.60
0.55
0.50
0.45
Macro-F1, test split
SFT (0.688)
one model call
two model calls
GAN final (0.555)
RAG (0.535)
Baseline (0.517)
Triple-Agent final (0.499)
Dual-Agent (0.467)
5
10
50
100
500
1000
Generated tokens per test instance (log scale)